

Rewriting a Deep Generative Model

David Bau, Steven Liu, Tongzhou Wang, Jun-Yan Zhu, Antonio Torralba
MIT CSAIL, Adobe

ECCV'20 Oral

Presented by Eungyeup Kim

Vision Seminar
01 SEP 2020

Motivation

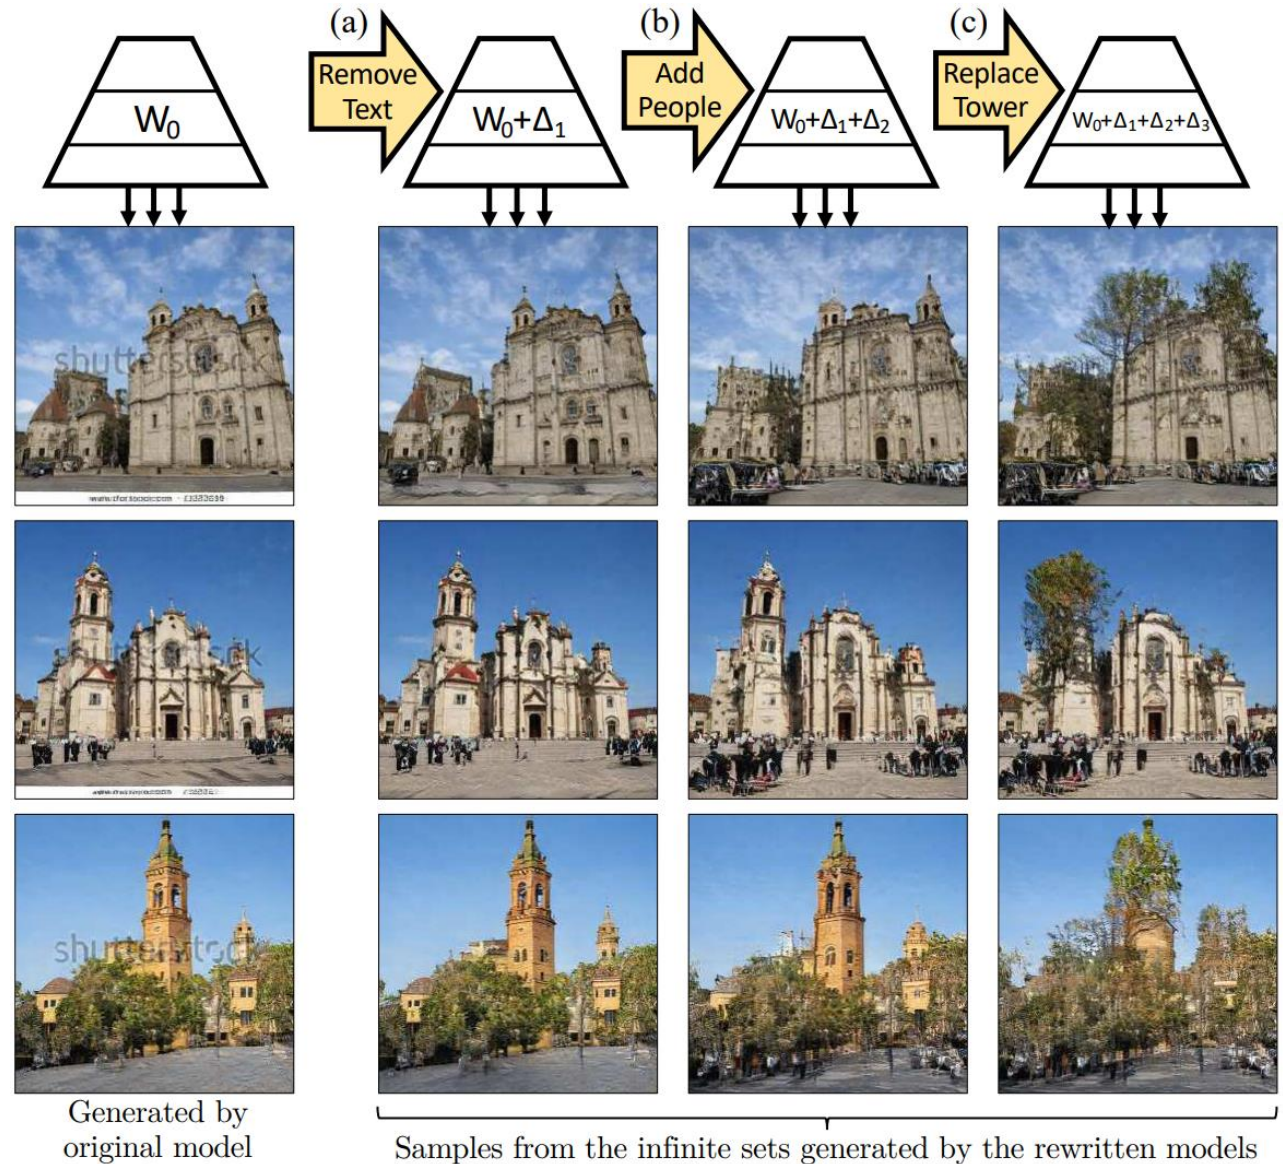
It has been obscure how the GANs encode the target distribution.

- A deep generative models such as GANs learn to model **a rich set of semantic and physical rules** about the target distribution.
 - However, it has been unclear how such rules are encoded in the network, or how a rule could be changed.
 - The ability to locate or edit those rules provides insights about what the model has captured and how the model can generalize to unseen scenarios.
 - We might want to create new data that never existed before. (Just like photoshop)
- ⇒ This paper presents the task of *model rewriting: add, remove, and alter the rules* of a pretrained deep network.

Introduction

Model Rewriting


- This method changes a set of desired rules by manipulating a layer of a deep network as a **linear associative memory**.
- By optimizing a layer with user-provided keys and values, the model can be **customized** without the large-scale of training time and computational cost.
- By adding, removing and changing the rules encoded in the model, totally unseen images can be generated by the user's intention.



Backgrounds

Viewing a Convolutional Layer as an Associative Memory

- In this paper, a convolutional layer is viewed as a memory in which many independent feature mappings are memorized.
- This is appropriate especially when the layer representations in the neighbors are **disentangled** from one another.



- Each key k denotes a single-location feature vector.
- A convolutional layer W stores an output value v corresponding to the k . **(b)**
- The same key will match many semantically similar locations across different images. **(c)**
- v renders shapes in a small region. **(d)**

Methods

Objective (Linear)

A traditional solution to update the model weights for generating desired images x_{*i} given z_i is as below:

$$\begin{aligned}\theta_1 &= \arg \min_{\theta} L_{smooth}(\theta) + \lambda L_{constraint}(\theta), \\ L_{smooth}(\theta) &= E_z[l(G(z; \theta_0), G(z; \theta))], \\ L_{constraint}(\theta) &= \sum_i l(x_{*i}, G(z_i; \theta)),\end{aligned}$$

where $G(z; \theta_0)$ represents a pre-trained generator, l a distance metric and x_{*i} an image containing desired changes. However, as the number of parameter θ is large, the generator quickly overfits to the new examples x_{*i} without good regularization. (..?)



$$\begin{aligned}W_1 &= \arg \min_W L_{smooth}(W) + \lambda L_{constraint}(W), \\ L_{smooth}(W) &= E_k[\|f(k; W_0) - f(k; W)\|^2], \\ L_{constraint}(W) &= \sum_i \|v_{*i} - f(k_{*i}; W)\|^2,\end{aligned}$$

Where given a layer L , f denotes a computation of L itself, k an input representation vector and v an output representation vector.

Methods

Objective (Linear)

As pretrained generator with layer L and its parameter W_0 ,

$$W_0 = \arg \min_W \sum_i \|v_i - Wk_i\|^2,$$
$$K = [k_1, k_2, \dots, k_i, \dots],$$
$$V = [v_1, v_2, \dots, v_i, \dots].$$

By using normal equation, we can obtain W_0 as below:

$$W_0 K K^T = V K^T$$

Suppose we wish to overwrite a single key to assign a new value $k_* \rightarrow v_*$ provided by the user. Then, a new parameter W_1 can be achieved as below:

$$W_1 = \arg \min_W \|V - WK\|^2,$$
$$\text{subject to } v_* = W_1 k_*.$$

By solving this constrained linear least-squares, we can obtain W_1 as below:

$$W_1 K K^T = W_0 K K^T + \Lambda k_*^T,$$
$$W_1 = W_0 + \Lambda (C^{-1} k_*)^T,$$

Where $C = K K^T$ (covariance).

Methods

Objective (Linear)

$$\begin{aligned}W_1 K K^T &= W_0 K K^T + \Lambda k_*^T, \\W_1 &= W_0 + \Lambda (C^{-1} k_*)^T.\end{aligned}$$

The update $\Lambda (C^{-1} k_*)^T$ is a rank-one matrix with rows all multiples of the vector $(C^{-1} k_*)^T$ or d^T .

1. The update is done in a particular straight-line direction $(C^{-1} k_*)^T$.
2. The update direction is determined only by the overall key statistics (covariance) and the specific targeted key k_* .
3. Only Λ depends on the target value v_* .

We can obtain the W_1 in a closed form.

Objective (Non-Linear)

However, what if f is a non-linear function, which is very common in our neural network setting..?

$$\Lambda_1 = \arg \min_{\Lambda} \sum_i \|v_* - f(k_*; W_0 + \Lambda d^T)\|^2,$$

Therefore, we can update the layer weights using $W_1 = W_0 + \Lambda_1 d^T$.

Experiments

User Interface

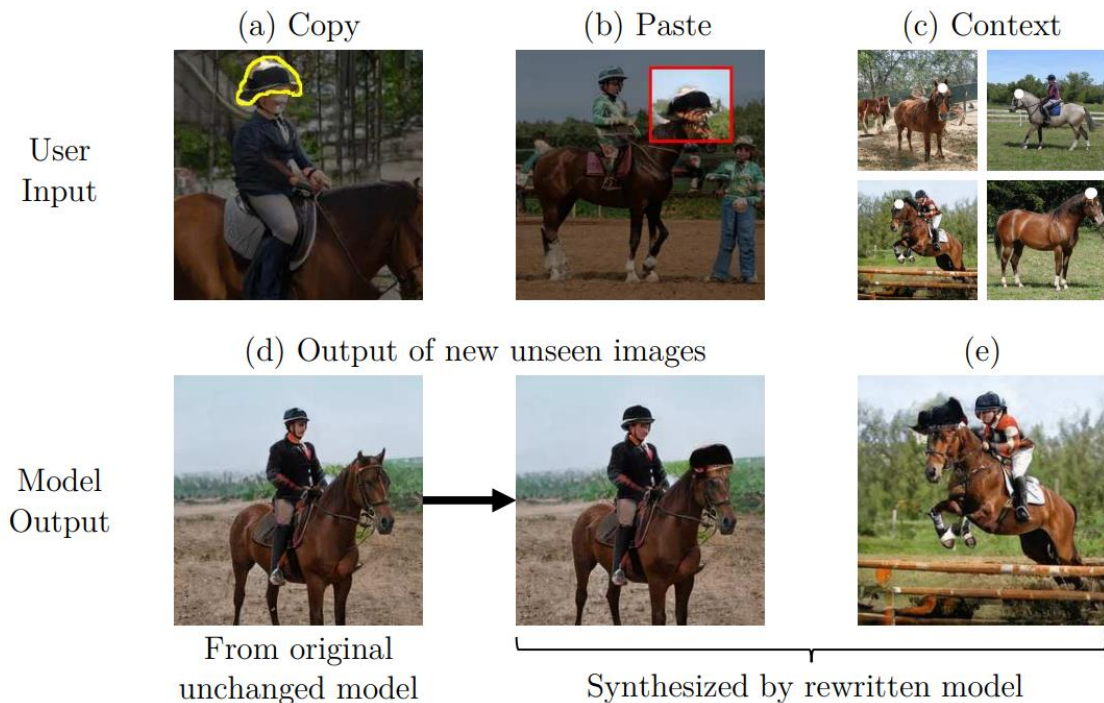


Fig. 3: The *Copy-Paste-Context* interface for rewriting a model. (a) **Copy:** the user uses a brush to select a region containing an interesting object or shape, defining the target value V_* . (b) **Paste:** The user positions and pastes the copied object into a single target image. This specifies the $K_* \rightarrow V_*$ pair constraint. (c) **Context:** To control generalization, the user selects target regions in several images. This establishes the updated direction d for the associative memory. (d) The edit is applied to the model, not a specific image, so newly generated images will always have hats on top of horse heads. (e) The change has generalized to a variety of different types of horses and poses (see more in Appendix A).

Experiments

Putting objects into a new context



Experiments

Removing undesired features


(a) Generated by
unchanged model



(b) Dissection:
zeroing 30 units



(c) Dissection:
zeroing 60 units

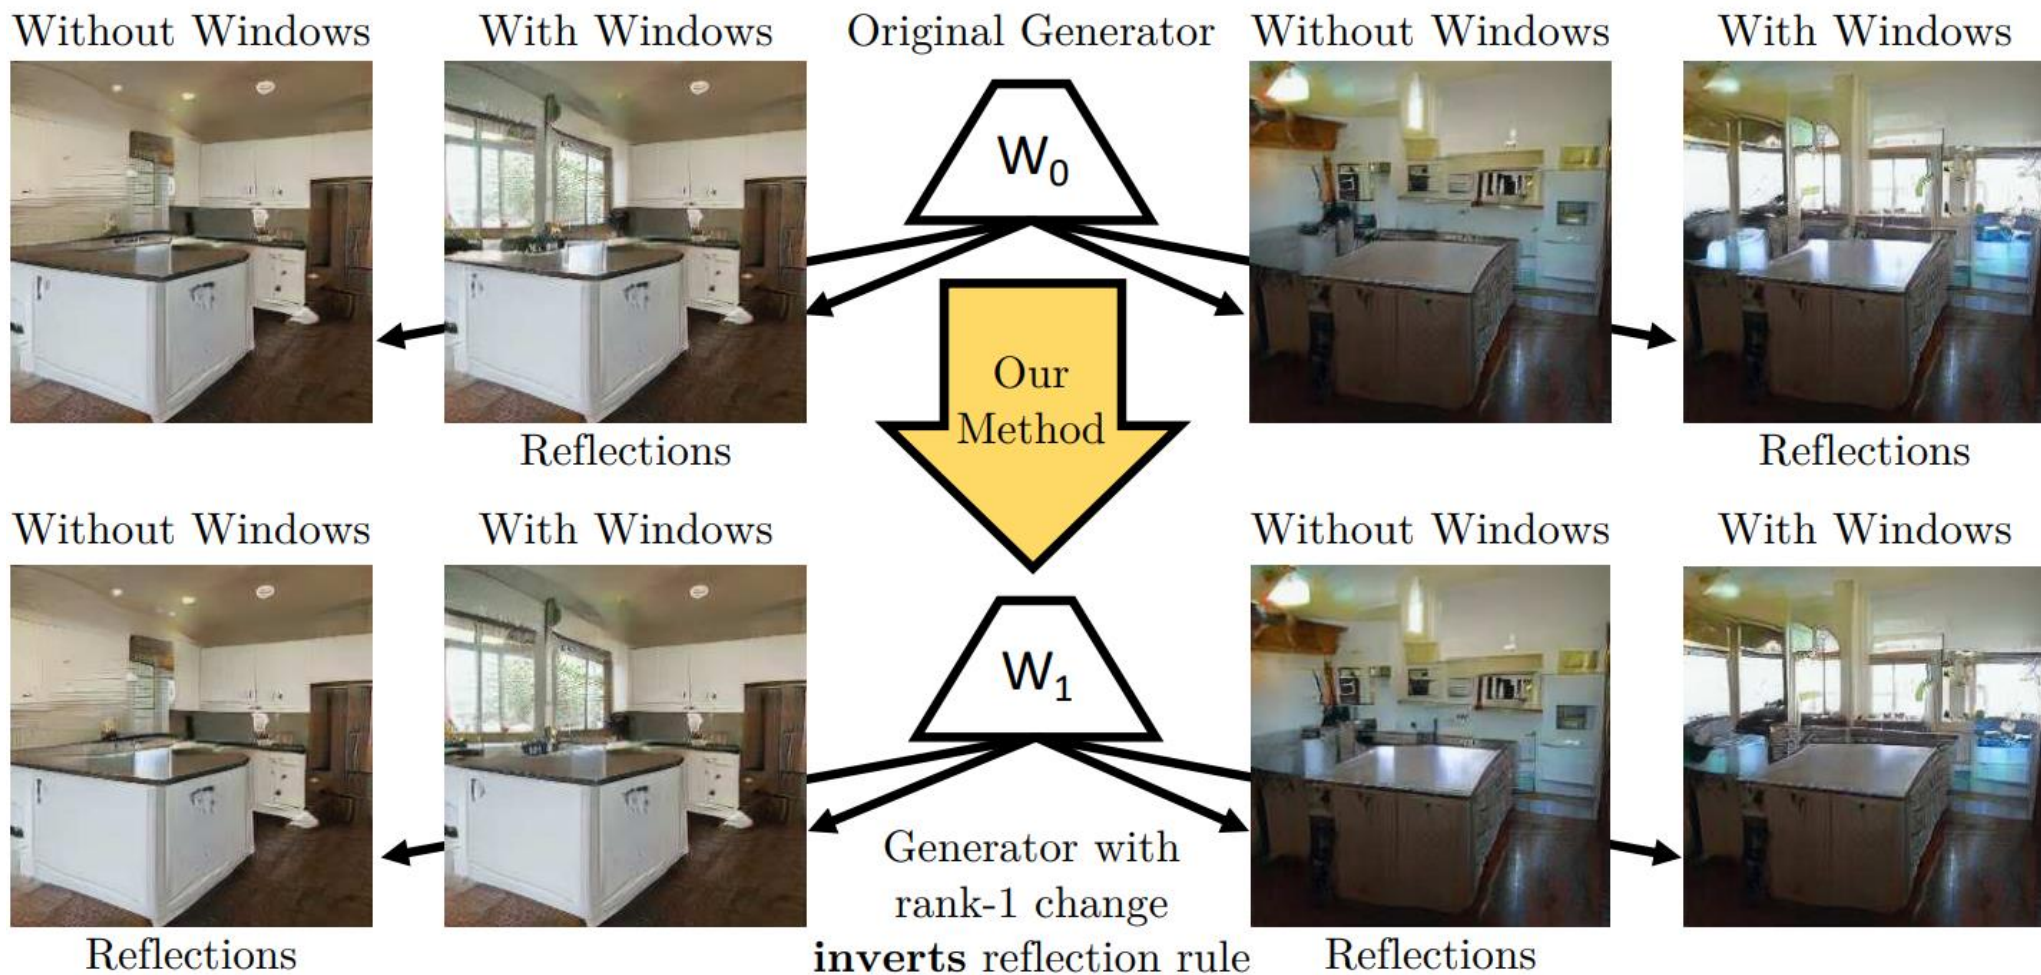


(d) Our method:
rank-1 update







Experiments

Changing contextual rules



Experiments

Which layer do we have to choose?



Thank you