

Scaling Local Self-Attention for Parameter Efficient Visual Backbones

CVPR '21 (*oral*)

(Uploaded to Arxiv on 23 Mar, 2021)

Presenter: Sungwon Hwang

May 3, 2021

Ashish Vaswani
Google Research

Prajit Ramachandran
Google Research

Aravind Srinivas
UC Berkeley

Niki Parmar
Google Research

Blake Hechtman
Google Research

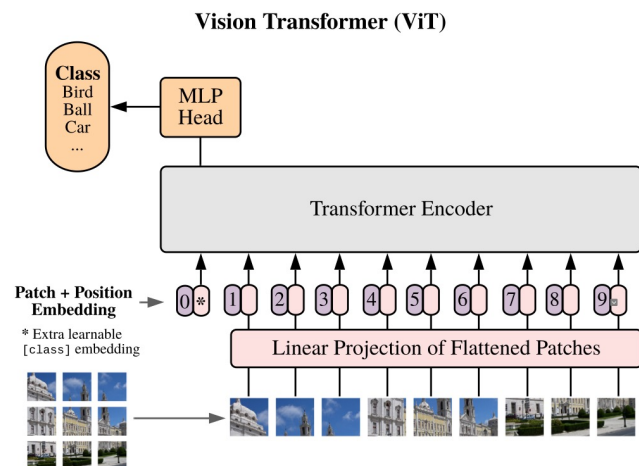
Jonathon Shlens
Google Research

Introduction

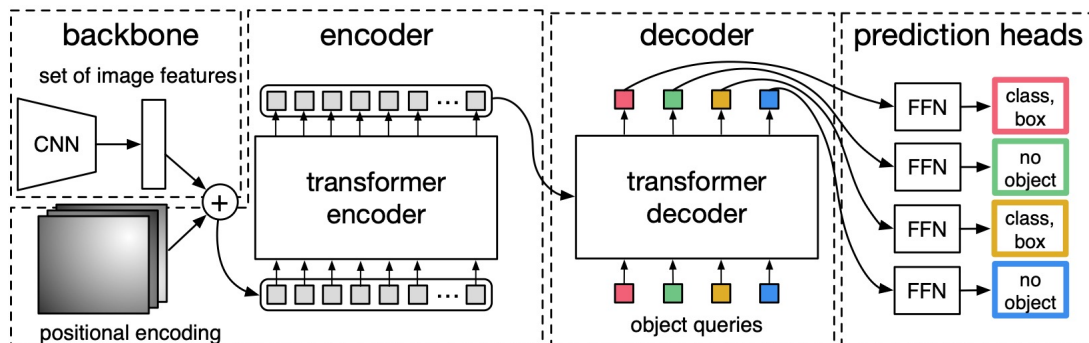
Advents of different non global self-attention for Vision

- **Global Attention**

- ViT (Vision Transformer)

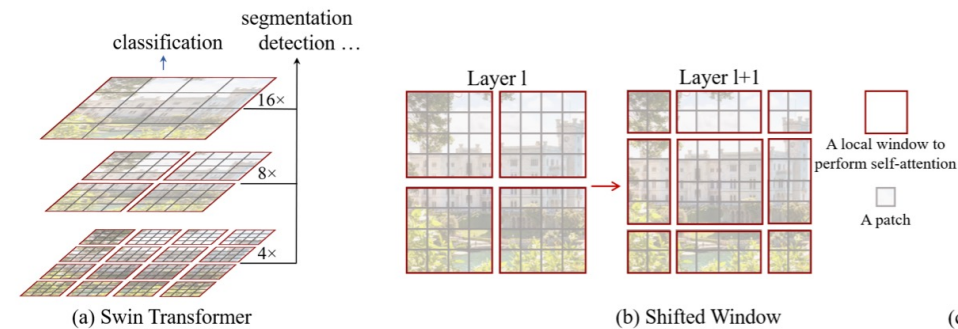


- DETR (Detection Transformer)



- **(Spatially) Local Attention**

- Swin (Sliding window) Transformer



- **(Selective) Local Attention**

- Deformable DETR

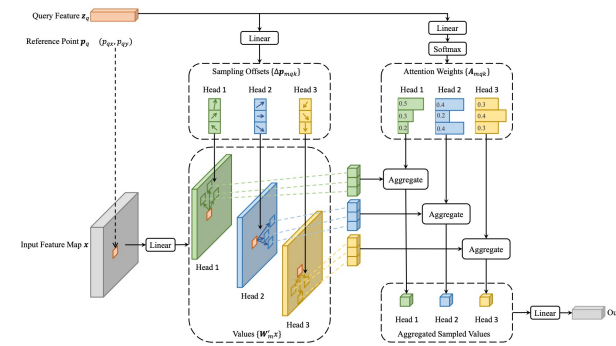


Figure 2: Illustration of the proposed deformable attention module.

Introduction

Advantages of self-attention over Convolution and why local self-attention for vision?

• Advantages of self-attention

1. **Content-based interactions** as opposed to content-independent interactions of convolutions
2. Parameter-independent scaling of receptive field size as opposed to parameter-dependent scaling of convolution
3. Empirical ability to **capture long-range dependencies** for use in larger images
4. Flexibility to handle and integrate multiple types of data
 - Pixels
 - Point Clouds
 - Any type of sequential information
 - Graphs

• Why local?

1. Nature of linguistic vs visual data?

- Language: unpredictable dependencies between words
 - > May be beneficial to attend global relationships
- Vision: Spatial locality of physical beings usually exists
 - > Is it THAT MUCH critical to understand global context in order to featurize local objects?
(ex. A picture of anyone's face can be taken anywhere)

2. Global attention requires high computational complexity.

- DETR Encoder: $O(H^2W^2C)$
 - > Quadratic growth of computation burden for a given increment of spatial size.

Methodology

Local self-attention as spatially varying convolutional filters

Design purpose: What if we can design self-attention as spatially-varying convolutional filters?

$$y_{ij} = \sum_{a,b \in \mathcal{N}(i,j)} f(i,j,a,b) x_{ab},$$

- **Convolutions**

$$f(i,j,a,b)^{conv} = W_{a-i,b-j}$$

- **Local Self-Attention**

$$\begin{aligned} f(i,j,a,b)^{self-att} &= \text{softmax}_{ab} \left(\underbrace{(W_Q x_{ij})^\top W_K x_{ab}}_{\text{Content - content interaction}} + \underbrace{(W_Q x_{ij})^\top r_{a-i,b-j}}_{\text{Content - geometry interaction}} \right) W_V \\ &= p_{a-i,b-j}^{ij} W_v \end{aligned}$$

Methodology

Local self-attention as spatially varying convolutional filters

$$y_{ij} = \sum_{a,b \in \mathcal{N}(i,j)} f(i,j,a,b) x_{ab},$$

- **Relative Spatial Attention**

$$\begin{aligned} f(i,j,a,b)^{self-att} &= \text{softmax}_{ab} \left((W_Q x_{ij})^\top W_K x_{ab} + \right. \\ &\quad \left. (W_Q x_{ij})^\top \boxed{r_{a-i,b-j}} W_V \right) \\ &\quad \text{Content - geometry} \\ &= p_{a-i,b-j}^{ij} W_v \quad \text{interaction} \end{aligned}$$

-1,-1	-1,0	-1,1	-1,2
0,-1	0,0	0,1	0,2
1,-1	1,0	1,1	1,2
2,-1	2,0	2,1	2,2


$$(a-i, b-j) \xrightarrow{\text{embed}} r_{a-i,b-j} = (r_{a-i}, r_{b-j}),$$
$$r_{a-i}, r_{b-j} \in \mathbb{R}^{d_{out}/2}$$

—> Relative Spatial Attention instead of positional encoding to get CLOSER to **TRANSLATIONAL EQUIVARIANCE**

(ViT does this positional embedding PATCHWISE => less translational equivariant => requires large-scale pretraining to have more access to larger samples of translations???)

Methodology

Blocked Local Self-Attention

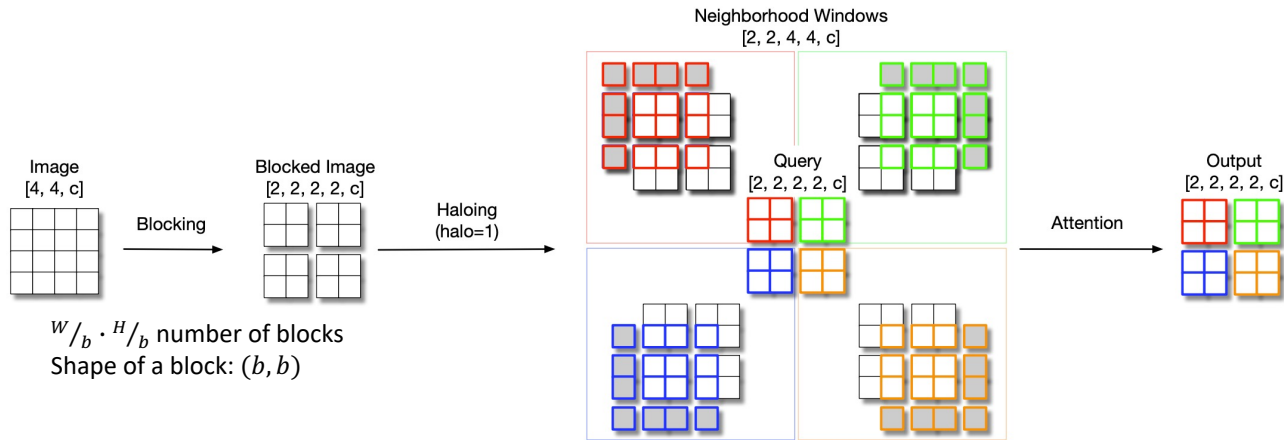


Method	Neighborhood Memory	Receptive Field	FLOPs Per Pixel	
Global	HWc	$HW \times HW$	$4(HW)^2c$	$\mathcal{O}((HW)^2)$
Per pixel windows	HWk^2c	$k \times k$	$4k^2c$	$\mathcal{O}(k^2)$
SASA [43]	$\frac{HW}{b^2}(b+2h)^2c$	$k \times k$, where $h = \lfloor \frac{k}{2} \rfloor$	$4(b+2h)^2c$	$\mathcal{O}((b+2h)^2)$
Blocked local (ours)	$\frac{HW}{b^2}(b+2h)^2c$	$(b+2h) \times (b+2h)$	$4(b+2h)^2c$	$\mathcal{O}((b+2h)^2)$

- Two straightforward methods are at the end of the spectrum
 - Global (No local window, or $k = HW$): **Quadratic computation** w.r.t spatial size
 - Per-pixel windows (The most similar to conventional convolutions): **Quadratic increment of neighborhood memory size** w.r.t window size
- Intuition for solution:** Adjacent pixels share most neighbors. ($k \times (k - 1)$ shared neighbors for $k \times k$ window)

Methodology

Blocked Local Self-Attention



Method	Neighborhood Memory	Receptive Field	FLOPs Per Pixel	
Global	HWc	$HW \times HW$	$4(HW)^2c$	$\mathcal{O}((HW)^2)$
Per pixel windows	HWk^2c	$k \times k$	$4k^2c$	$\mathcal{O}(k^2)$
SASA [43]	$\frac{HW}{b^2}(b+2h)^2c$	$k \times k$, where $h = \lfloor \frac{k}{2} \rfloor$	$4(b+2h)^2c$	$\mathcal{O}((b+2h)^2)$
Blocked local (ours)	$\frac{HW}{b^2}(b+2h)^2c$	$(b+2h) \times (b+2h)$	$4(b+2h)^2c$	$\mathcal{O}((b+2h)^2)$

- **Solution:** Blocks can safely share the same neighbors: **Pixels in blocks attend to shared, haloed neighbors**

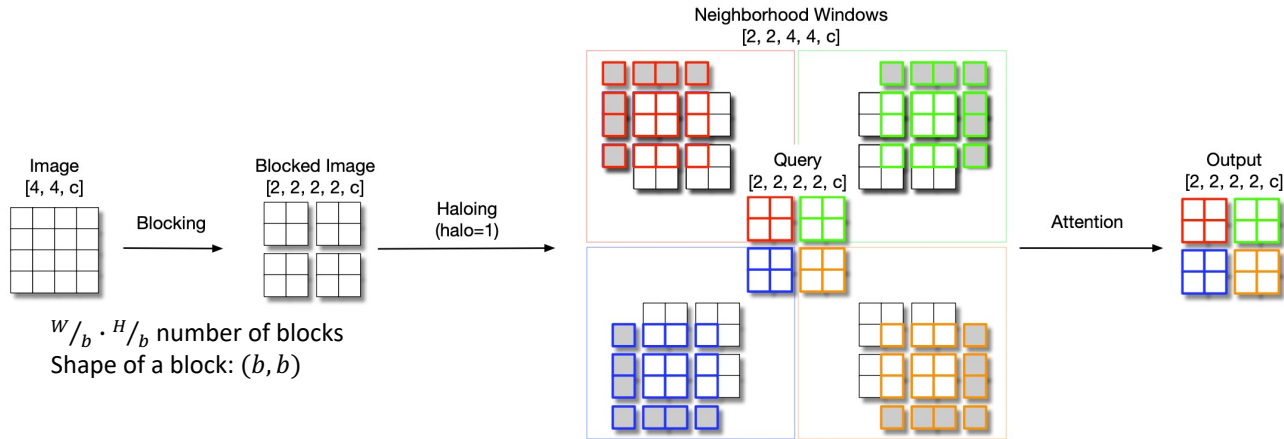
- Query: Pixels in **blocks** -> $(W/b \cdot H/b, c)$
- Key, value: Pixels in **halos** -> $((b+2h)^2, c)$

-> Manageable neighborhood memory while controlling FLOPs per pixel

-> Do NOT need to have larger number of parameters for larger receptive field compared to CNN.

Methodology

Blocked Local Self-Attention



Method	Neighborhood Memory	Receptive Field	FLOPs Per Pixel	
Global	HWc	$HW \times HW$	$4(HW)^2c$	$\mathcal{O}((HW)^2)$
Per pixel windows	HWk^2c	$k \times k$	$4k^2c$	$\mathcal{O}(k^2)$
SASA [43]	$\frac{HW}{b^2}(b+2h)^2c$	$k \times k$, where $h = \lfloor \frac{k}{2} \rfloor$	$4(b+2h)^2c$	$\mathcal{O}((b+2h)^2)$
Blocked local (ours)	$\frac{HW}{b^2}(b+2h)^2c$	$(b+2h) \times (b+2h)$	$4(b+2h)^2c$	$\mathcal{O}((b+2h)^2)$

- Comparisons of Neighborhood Memory & FLOPs: per-pixel windows vs. Blocked local

- Set equal receptive field for fair comparison: $k = b + 2h$

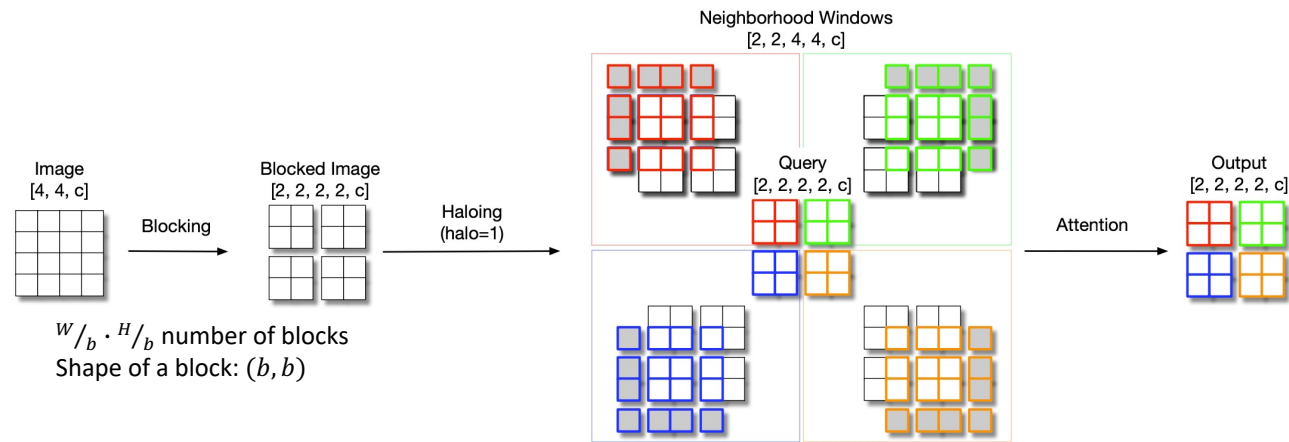
- Neighborhood Memory: $\frac{\text{Per-pixel}}{\text{Blocked local}} = b^2$

- FLOPs per pixel: $\frac{\text{Per-pixel}}{\text{Blocked local}} = 1$

- => Manageable neighborhood memory while keeping FLOPs per pixel


Methodology

Blocked Local Self-Attention vs. SASA: Larger receptive fields are more important than translational equivariance



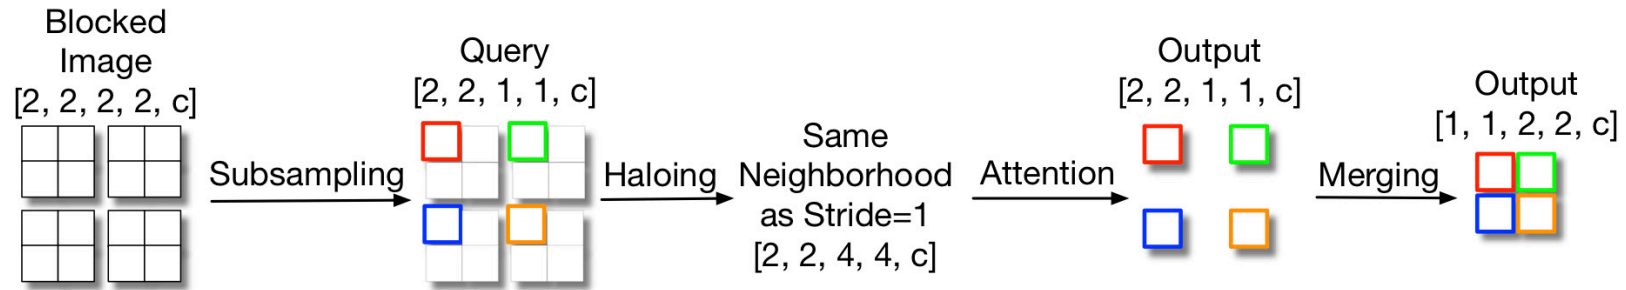
Method	Neighborhood Memory	Receptive Field	FLOPs Per Pixel	
Global	HWc	$HW \times HW$	$4(HW)^2c$	$O((HW)^2)$
Per pixel windows	HWk^2c	$k \times k$	$4k^2c$	$O(k^2)$
SASA [43]	$\frac{HW}{b^2}(b+2h)^2c$	$k \times k$, where $h = \lfloor \frac{k}{2} \rfloor$	$4(b+2h)^2c$	$O((b+2h)^2)$
Blocked local (ours)	$\frac{HW}{b^2}(b+2h)^2c$	$(b+2h) \times (b+2h)$	$4(b+2h)^2c$	$O((b+2h)^2)$

- SASA: gives up receptive field for translational equivariance



Methodology

Down-sampling with local self-attention



- Just like strided convolution with stride of b

Model Configuration

HaloNet: ResNet-like structure & empirical comparisons with Efficient Net

HaloNet

Output Resolution	Layers
$\frac{s}{4} \times \frac{s}{4}$	7×7 conv stride 2, 64 3×3 max pool stride 2
$\frac{s}{4} \times \frac{s}{4}$	$\left\{ \begin{array}{l} 1 \times 1, 64 \\ \text{attention}(b, h), 64 \cdot r_v \\ 1 \times 1, 64 \cdot r_b \end{array} \right\} \times 3$
$\frac{s}{8} \times \frac{s}{8}$	$\left\{ \begin{array}{l} 1 \times 1, 128 \\ \text{attention}(b, h), 128 \cdot r_v \\ 1 \times 1, 128 \cdot r_b \end{array} \right\} \times 3$
$\frac{s}{16} \times \frac{s}{16}$	$\left\{ \begin{array}{l} 1 \times 1, 256 \\ \text{attention}(b, h), 256 \cdot r_v \\ 1 \times 1, 256 \cdot r_b \end{array} \right\} \times l_3$
$\frac{s}{32} \times \frac{s}{32}$	$\left\{ \begin{array}{l} 1 \times 1, 512 \\ \text{attention}(b, h), 512 \cdot r_v \\ 1 \times 1, 512 \cdot r_b \end{array} \right\} \times 3$
$\frac{s}{32} \times \frac{s}{32}$	$1 \times 1, d_f$
1×1	global average pooling fc, 1000

Value 의 linear projection 시 dimension 을 늘려주면 됨

ResNet


layer name	output size	18-layer	34-layer	50-layer	101-layer	152-layer
conv1	112×112			7×7, 64, stride 2		
				3×3 max pool, stride 2		
conv2_x	56×56	$\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$
conv3_x	28×28	$\begin{bmatrix} 3 \times 3, 128 \\ 3 \times 3, 128 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 128 \\ 3 \times 3, 128 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 8$
conv4_x	14×14	$\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 23$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 36$
conv5_x	7×7	$\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$
	1×1	average pool, 1000-d fc, softmax				
FLOPs		1.8×10^9	3.6×10^9	3.8×10^9	7.6×10^9	11.3×10^9

Table 2. HaloNet model family specification.

Experiments

Comparison to Conv Exp #1) ImageNet classification accuracy vs. EfficientNet

- HaloNet can be **trained over ImageNet from scratch** & yield state-of-the-art performance
 - ViT / BiT: Requires pre-training on larger datasets
 - Imagenet-21k
 - JFT-300M



HaloNet Model	b	h	r_v	r_b	Total Layers	l_3	s	d_f	Params (M)	EfficientNet Params (M)	EfficientNet Image Size (M)
H0	8	3	1.0	0.5	50	7	256	–	5.5	B0: 5.3	224
H1	8	3	1.0	1.0	59	10	256	–	8.1	B1: 7.8	240
H2	8	3	1.0	1.25	62	11	256	–	9.4	B2: 9.2	260
H3	10	3	1.0	1.5	65	12	320	1024	12.3	B3: 12	300
H4	12	2	1.0	3	65	12	384	1280	19.1	B4: 19	380
H5	14	2	2.5	2	98	23	448	1536	30.7	B5: 30	456
H6	8	4	3	2.75	101	24	512	1536	43.4	B6: 43	528
H7	10	3	4	3.5	107	26	600	2048	67	B7: 66	600

Experiments

Comparison to Conv Exp #2) Transfer of convolutional components to self-attention

- Architecture variation: Squeeze and Excitation (SE) / channel attention after spatial convolution (SiLU/Swish-1)
- Regularization: Random Augment (RA) / Label Smoothing (LS)
- HaloNet leverages from regularization more than ResNet-50.
 - **Label Smoothing(LS) + RandomAugment (RA)** yields 1.3% increase for Halo / 0.8 for ResNet-50
 - Usually larger convolutional models benefit from regularization -> HIGHER EXPRESSIBILITY OF HALONET?
- Components that doesn't affect HaloNet
 - **Squeeze and Excitation (SE)**
 - Self-attention module already has higher DOF of expressibility than channel-wise attention block (SE).

Components	HaloNet Accuracy	Baseline Δ	ResNet Accuracy	Baseline Δ
Baseline	78.6	0.0	77.6	0.0
+ LS	79.7	1.1	78.1	0.5
+ LS, RA	79.9	1.3	78.4	0.8
+ SE	78.6	0.0	78.6	1.0
+ SE, SiLU/Sw1	79.0	0.4	78.9	1.3
+ LS, SE	79.7	1.1	78.9	1.3
+ LS, SE, SiLU/Sw1	79.9	1.3	79.1	1.5
+ LS, SE, SiLU/Sw1, RA	80.5	1.9	79.5	1.9

Experiments

Comparison to Conv Exp #3) Comparative study on Object Detection / Instance Segmentation (COCO)

- Last 3 layers of ResNet baselines are added with local attention block to compare with ResNet
- Results to note:
 - Still suffering from localizing small objects
 - Feels like translational equivariance is critical for fine-grained localization. Maybe try SASA (local attention mask)?
 - Larger box size perform better for localizing smaller objects?

Model	AP^{bb}	AP_s^{bb}	AP_m^{bb}	AP_l^{bb}	AP^{mk}	AP_s^{mk}	AP_m^{mk}	AP_l^{mk}	Speed (ms)	Train time (hrs)
R50 baseline in lit	42.1	22.5	44.8	59.1	37.7	18.3	40.5	54.9	409	14.6
R50 + SE (our baseline)	44.5 (+2.4)	25.5	47.7	61.2	39.6 (+1.9)	20.4	42.6	57.6	446	15.2
R50 + SE + Local Att ($b = 8$)	45.2 (++)	25.4	48.1	63.3	40.3 (++)	20.5	43.1	59.0	540	15.8
R50 + SE + Local Att ($b = 32$)	45.4 (++)	25.9	48.2	63.0	40.5 (++)	21.2	43.5	58.8	613	16.5
R101 + SE (our baseline)	45.9 (+3.8)	25.8	49.5	62.9	40.6 (+2.9)	20.9	43.7	58.7	740	17.9
R101 + SE + Local Att ($b = 8$)	46.8 (++)	26.3	50.0	64.5	41.2 (++)	21.4	44.3	59.8	799	18.4

Table 5. **Accuracies on object detection and instance segmentation.** We experiment with two settings for self-attention in the last stage: A block size of (b) of 8 and a halo size (h) of 3 and also with ($b = 32, h = 3$) for ResNet-50. bb (bounding box) refers to detection, and mk (mask) refers to segmentation. The identifiers s , m , and l refer to small, medium, and large objects respectively. Speed is measured as the milliseconds taken by only the backbone (and not the FPN) for a batch size of 32 on 2 TPUv3 cores. The train time the total training time calculated from the peak images/sec of the Mask-RCNN training run on 8 TPUv3 cores with a batch size of 64.

Experiments

Comparison to Transformers for Vision

- Pretrain: Public ImageNet-21k
- Finetune: ImageNet.
- *4×4 patch*: Replacing the convolutional stem with linear projection of *4×4 patch* (Just like ViT)
- *Conv-12*: First 2 stages are made with convolutions
- **Good parameter trade-off / inference speed trade-off for given accuracy**

Conv Stages	Attention Stages	Top-1 Acc (%)	Norm. Train Time
-	1, 2, 3, 4	84.9	1.9
1	2, 3, 4	84.6	1.4
1, 2	3, 4	84.7	1.0
1, 2, 3	4	83.8	0.5

Table 4. Replacing attention layers with convolutions in stages 1 and 2 exhibit the best speed vs. accuracy tradeoff. All the models had about 67 million parameters and the train and inference times are normalized to the corresponding times for EfficientNet B7. Please see Figure 8 for a comparison of step time with other HaloNet models.


Model	Parameters (Millions)	Pretraining Image Size (Pixels)	Pretraining Step Time (32 per core)	Finetuning Image Size	Finetuning Top-1 Accuracy (%)	Inference Speed img/sec/core
H4 (base 128)	85	256	377 ms	384/512	85.6/85.8	121.3/48.6
H4 (base 128, 4 × 4 patch)	85	256	366 ms	384/512	85.4/85.4	125.7/56.5
H4 (base 128, Conv-12)	87	256	213 ms	384/512	85.5/85.8	257.6/120.2
ViT-L/16	300	224	445 ms	384/512	85.2/85.3	74.6/27.4
BiT-M	928	224	1021 ms	384	85.4	54.2

Table 6. HaloNet models pretrained on ImageNet-21k perform well when finetuned on ImageNet. For HaloNet and ViT, we finetuned on 384×384 and 512×512 size images. The pretraining step time reports the TPUv3 compute time for a batch size of 32 per core. The inference speed is also computed on a single TPUv3 core.

Experiments

HaloNet model study #1) HaloNet vs. SASA: Is it okay to relax translational equivariance for larger receptive field?

- $b = 8, h = 3, k = 7$



Experiments

HaloNet model study #2) Size of receptive field

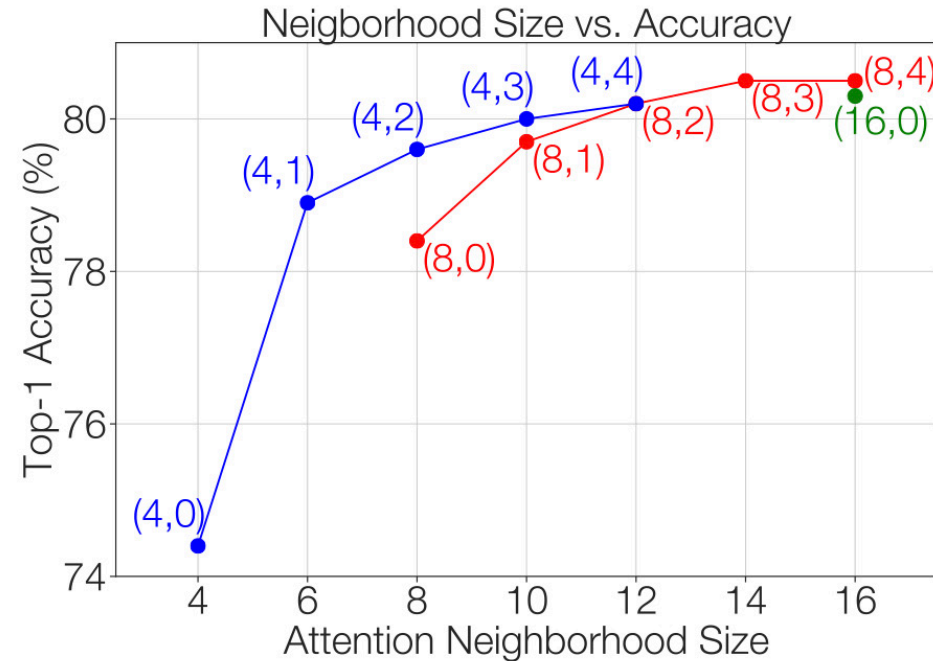


Figure 7. **Increasing window sizes improves accuracy up to a point.** The experiments in the graph have been annotated with their block size (b), halo size (h), $h = 0$ implies attention with *non-overlapping* blocks