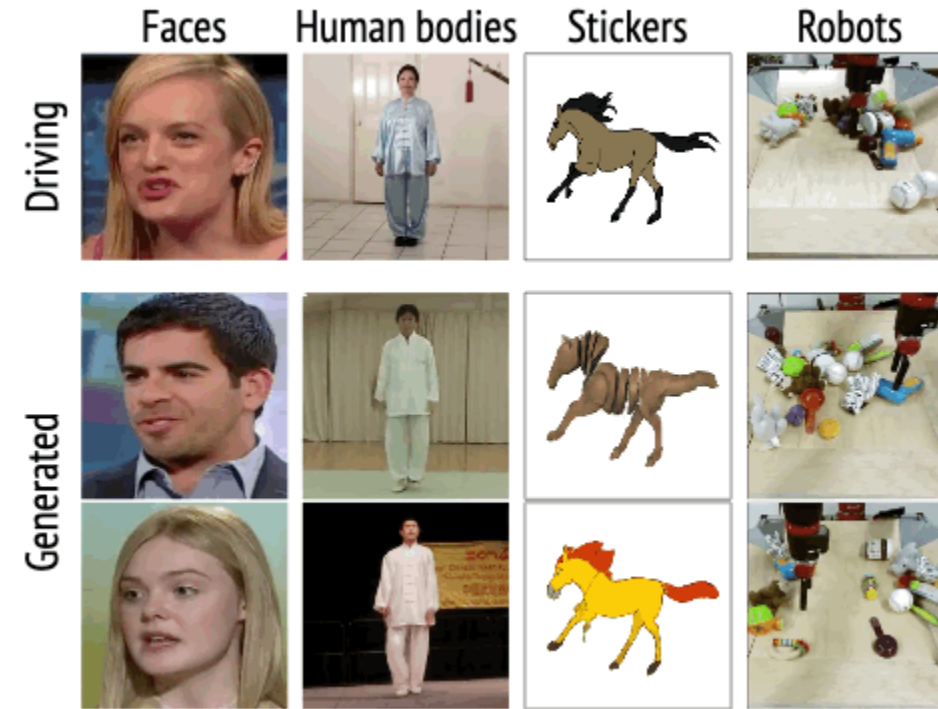


First Order Motion Model for Image Animation

Aliaksandr Siarohin, Stéphane Lathuilière, Sergey Tulyakov, Elisa Ricci, Nicu Sebe

Kihong Kim

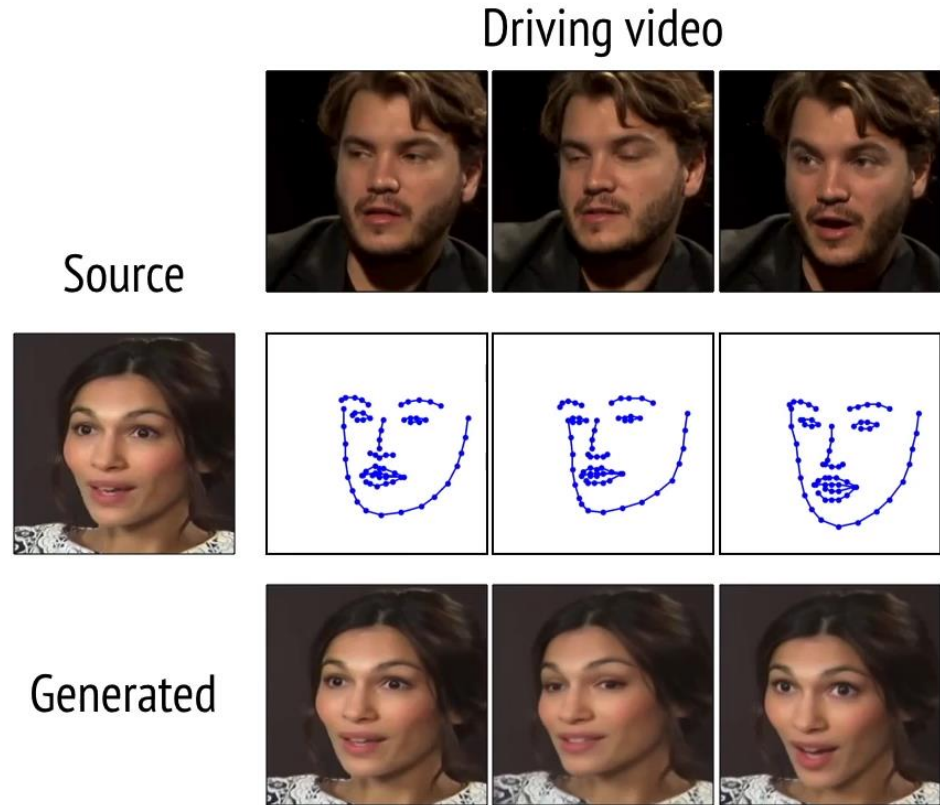
Introduction




- Image animation
 - Animating object in a source image according to the motion of a driving video

Background

- Prior works
 - Are object specific
 - Require landmark detectors
 - Impose too strong motion prior
- Our Work
 - Object-agnostic model
 - Does not require object-specific prior
 - Animates multiple objects categories




Self-supervised Training




- Direct supervision is not available
 - Train Keypoint Detector in a self-supervised manner
- Extracting each source & target frame from the same video
 - Reconstruct the training videos by combining a single frame and a learned latent motion

Overview



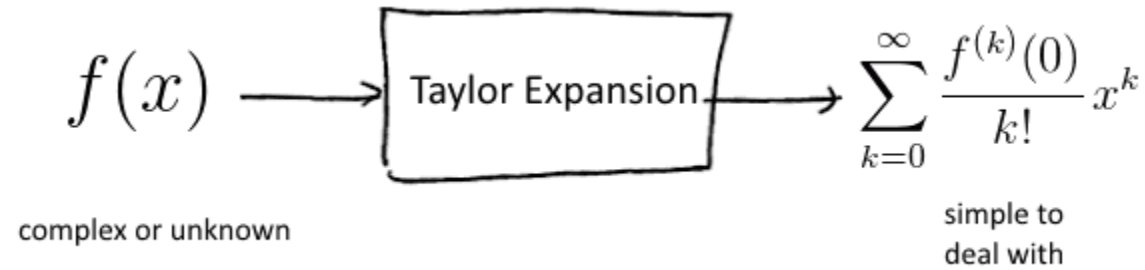
- Motion Module $\vec{T}_{S \leftarrow D}$
 - Predict a dense motion field from a source and driving frame ($T_{S \leftarrow D}$)
- Image Generation Module
 - Renders an image of the source object moving as provided in the driving video

Motion Module – Keypoint Detector




- Keypoint Detector predict **keypoint displacement** and **local affine transformation**
- Approximate $T_{S \leftarrow D}$ by its first order Taylor expansion in a neighborhood of the keypoint locations
- $$T_{S \leftarrow D}(p) = T_{S \leftarrow D}(p_k) + \left(\frac{d}{dp} T_{S \leftarrow D}(p) \Big|_{p=p_k} \right)$$


Taylor expansion



$$\begin{aligned} e^x &= f(0)\frac{x^0}{0!} + f'(0)\frac{x^1}{1!} + f''(0)\frac{x^2}{2!} + f'''(0)\frac{x^3}{3!} + f^{(4)}(0)\frac{x^4}{4!} + f^{(5)}(0)\frac{x^5}{5!} + \dots \\ &= \frac{x^0}{0!} + \frac{x^1}{1!} + \frac{x^2}{2!} + \frac{x^3}{3!} + \frac{x^4}{4!} + \frac{x^5}{5!} + \dots \\ &= \sum_{n=0}^{\infty} \frac{x^n}{n!} \end{aligned}$$




Motion Module – Keypoint Detector




- $T_{X \leftarrow R}(p) = T_{X \leftarrow R}(p_k) + \left(\frac{d}{dp} T_{X \leftarrow R}(p) \Big|_{p=p_k} \right)$
- Assume an abstract reference frame **R**
 - $T_{S \leftarrow D} = T_{S \leftarrow R} \circ T_{R \leftarrow D} = T_{S \leftarrow R} \circ T_{D \leftarrow R}^{-1}$
 - **R** allow us to independently process **S** and **D**

Motion Module – Keypoint Detector




- $T_{S \leftarrow D} = T_{S \leftarrow R} \circ T_{R \leftarrow D} = \mathbf{T}_{S \leftarrow R} \circ \mathbf{T}_{D \leftarrow R}^{-1}$
- $T_{X \leftarrow R}(p) \cong \left\{ T_{X \leftarrow R}(p_k), \frac{d}{dp} T_{X \leftarrow R}(p)|_{p=p_1} \right\} + \dots + \left\{ T_{X \leftarrow R}(p_k), \frac{d}{dp} T_{X \leftarrow R}(p)|_{p=p_k} \right\}$
- $T_{S \leftarrow D}(z) \approx T_{S \leftarrow R}(p_k) + J_k(z - T_{D \leftarrow R}(p_k))$
 - $J_k = \left(\frac{d}{dp} T_{S \leftarrow R}(p)|_{p=p_k} \right) \left(\frac{d}{dp} T_{D \leftarrow R}(p)|_{p=p_k} \right)^{-1}$

Motion Module - Dense Motion



- $T_{S \leftarrow D(z)} \approx T_{S \leftarrow R(p_k)} + J_k(z - T_{D \leftarrow R(p_k)})$
- Dense Motion network combines the local approximations to obtain dense motion field $\hat{T}_{S \leftarrow D(z)}$
- Dense Motion network outputs an occlusion mask $\hat{O}_{S \leftarrow D}$

Motion Module - Dense Motion

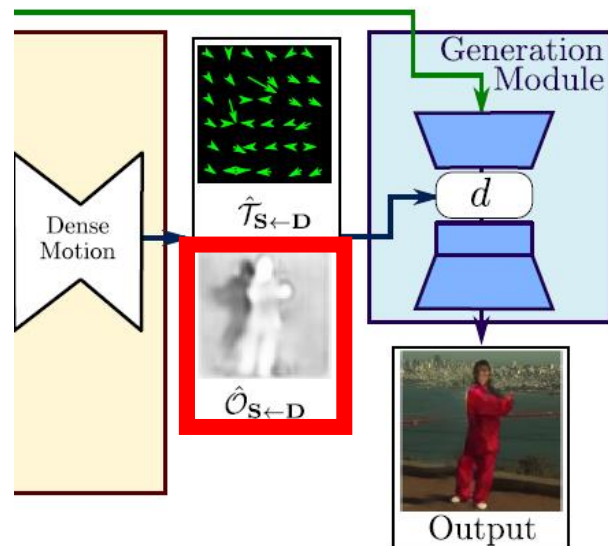


- $\hat{T}_{S \leftarrow D(Z)} = M_0 Z + \sum_{k=1}^K M_k (T_{S \leftarrow R(p_k)} + J_k(Z - T_{D \leftarrow R(p_k)}))$

- $M_k(z) = \exp\left(\frac{(T_{D \leftarrow R}(p_k) - z)^2}{\sigma}\right) - \exp\left(\frac{(T_{S \leftarrow R}(p_k) - z)^2}{\sigma}\right)$


- Heatmap indicate to the dense motion network where each transformation happens

Occlusion-Aware Image Generation



- $\xi' = \hat{o}_{S \leftarrow D} \odot f_w(\xi, \hat{\tau}_{S \leftarrow D})$
- Occluded parts in S cannot be recovered by image-warping and thus should be inpainted
- Mask out feature map regions that should be inpainted

Generation Module



- Generation Module renders an image of the source object moving as provided in the driving video
- warps the source image according to $\hat{\mathcal{T}}_{S \leftarrow D}$ and inpaints the image parts that are occluded in the source image

Loss function

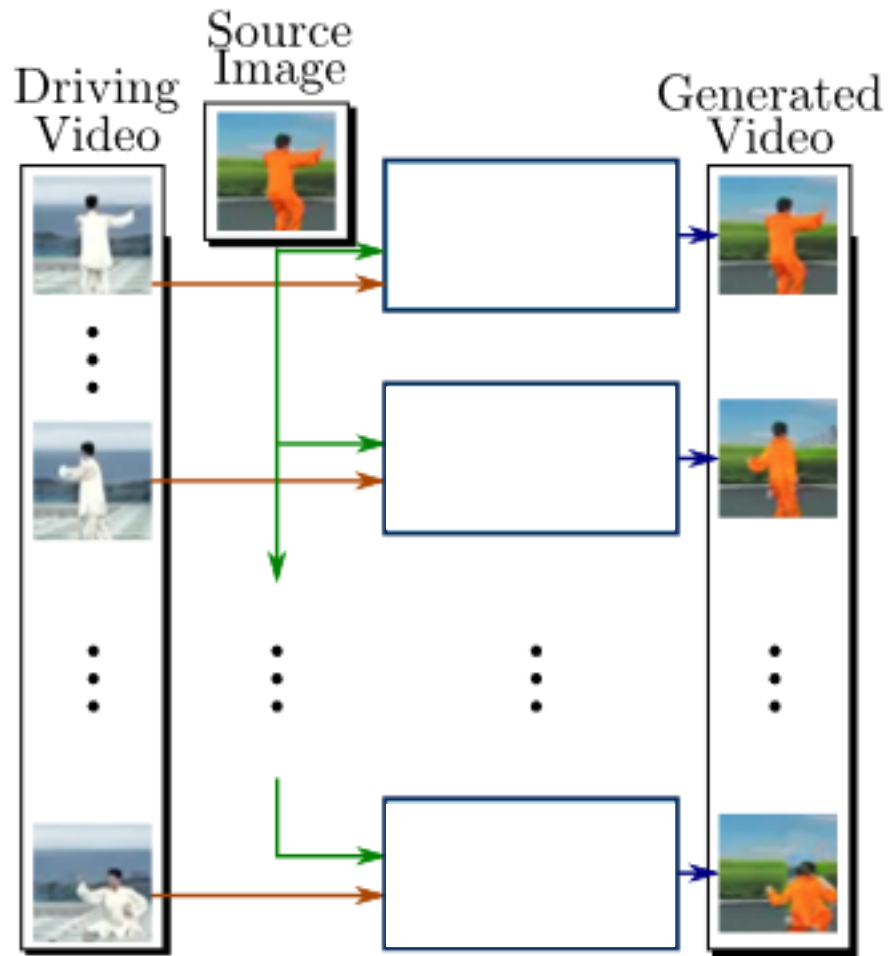
- **Reconstruction Loss**

- $L_{rec}(\hat{D}, D) = \sum_{i=1}^I |N_i(\hat{D}) - N_i(D)|$
- Perceptual loss using the pre-trained VGG-19
- Multiple resolution
 - 256x256, 128x128, 64x64, 32x32

- **Equivariance Constraint**

- Our keypoint predictor doesn't require any keypoint annotations during training.
- This may lead to unstable performance
- $T_{X \leftarrow R}(p_k) \equiv T_{X \leftarrow Y} \circ T_{Y \leftarrow R}(p_k)$

Testing Stage – Relative Motion Transfer



- Animation using absolute coordinates



- Animation using relative coordinates



Experiment



- **VoxCeleb**
cropping face from image using bounding box
19,522 training videos and 525 test videos
- **UvA-Nemo**
Facial analysis dataset
1116 training videos and 124 test videos
- **BAIR robot pushing**
42,880 training and 128 test videos
30 frame long, 256x256 resolution
- **Tai-Chi-HD**
3,049 training, 285 testing
128 to 1024 frames

Ablation Study

Table 1: Quantitative ablation study for video reconstruction on *Tai-Chi-HD*.

	\mathcal{L}_1	<i>Tai-Chi-HD</i> (AKD, MKR)	AED
<i>Baseline</i>	0.073	(8.945, 0.099)	0.235
<i>Pyr.</i>	0.069	(9.407, 0.065)	0.213
<i>Pyr.+$\mathcal{O}_{S \leftarrow D}$</i>	0.069	(8.773, 0.050)	0.205
<i>Jac. w/o Eq. (12)</i>	0.073	(9.887, 0.052)	0.220
<i>Full</i>	0.063	(6.862, 0.036)	0.179

- L1 Distance
- AKD (Average Keypoint Distance)
 - Evaluate whether the motion of the input video is preserved
- MKR (Missing Keypoint Rate)
 - Evaluate the appearance quality of each generated frame
- AED (Average Euclidean Distance)
 - Evaluate Euclidean distance between G.T and generated frame representation




Figure 3: Qualitative ablation on *Tai-Chi-HD*.

Ablation Study



Table 3: Video reconstruction: comparison with the state of the art on four different datasets.

	\mathcal{L}_1	<i>Tai-Chi-HD</i> (AKD, MKR)	AED	\mathcal{L}_1	<i>VoxCeleb</i> AKD	AED	\mathcal{L}_1	<i>Nemo</i> AKD	AED	<i>Bair</i> \mathcal{L}_1
X2Face [41]	0.080	(17.654, 0.109)	0.272	0.078	7.687	0.405	0.031	3.539	0.221	0.065
Monkey-Net [29]	0.077	(10.798, 0.059)	0.228	0.049	1.878	0.199	0.018	1.285	0.077	0.034
Ours	0.063	(6.862, 0.036)	0.179	0.043	1.294	0.140	0.016	1.119	0.048	0.027

- Our approach is able to generate significantly better looking videos in which each body part is independently animated

Conclusion

- Mathematical formulation describes the motion field
 - a set of keypoints displacement and local affine transformations
- Dense Motion Network produce Occlusion mask to inpaint occluded region
 - Occluded parts in S cannot be recovered by image-warping
- Test our method on four different datasets containing various objects
 - Our approach outperforms existing method

Thank You !